

2-3 Threads

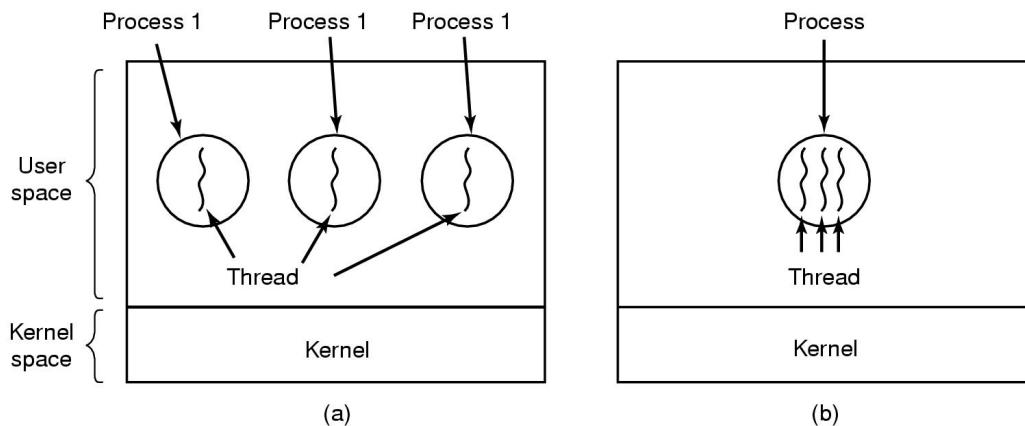
Stallings: 4.1, 4.2, 4.3, 4.6

Introduction to Threads

- The overhead of using multiple “full-blown” processes can limit the performance of many classes of applications (e.g., server applications)
- The concept of threads, a.k.a. **lightweight processes**, came under study in the 1980s.
- Now, many commercial and public-domain thread packages are available and are widely used.
- Processor architectures also support threads
- Multi-threaded programming is often challenging
- **In many thread packages, semantics is closely tied to the operating system...**

Definition

- Thread: execution of a sequence of instructions within an address space
- A process can contain one or more threads
- Multiple threads wandering through the same program **concurrently**, maybe simultaneously!

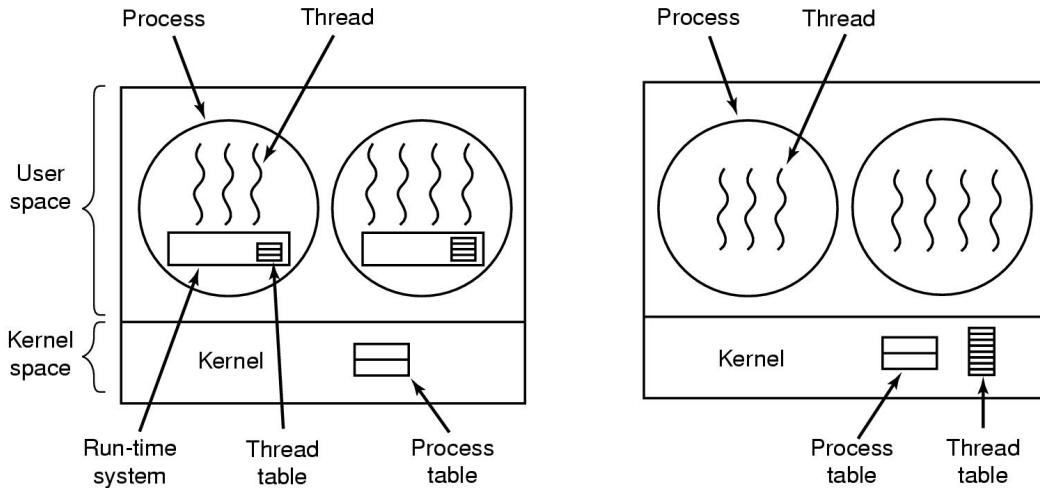


User-level vs. Kernel-level Threads

- Threads may, or may not, be supported by the operating system
- User-level threads
 - Thread support implemented within one user-level process
 - Part of the user-level program (run-time system) manages and schedules the threads, all within one process
 - Threads are **not** known to the kernel
- Kernel-level threads
 - Threads are **known to kernel**
 - State information (like PCB) maintained on thread basis
 - Usually can be **scheduled** independently

User-Level vs. Kernel-Level Support

- ULT: kernel sees only the **process** (one thread)
- KLT: kernel maintains info on each thread



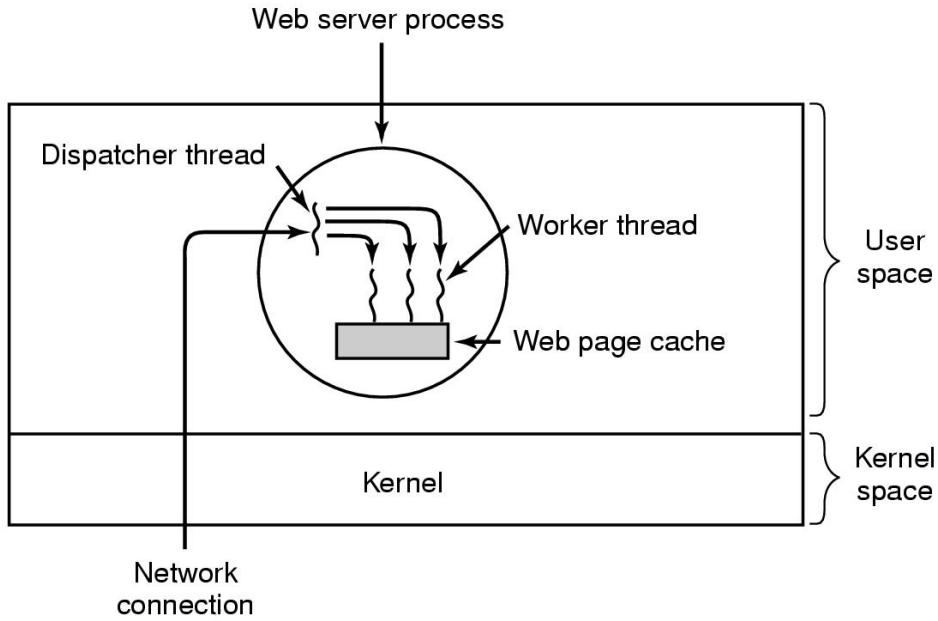
- Implications for concurrency? performance?

Multi-threaded programming

- Program is written as a set of threads
- The threads execute concurrently within single process address space.
- Threads are designed to cooperatively implement some functionality. E.g.,
 - web server
 - parallel processing application
 - parallel database application
- Threads must interact via synchronization primitives
 - E.g., semaphores, mutexes, condition variables, etc.
 - Otherwise? inconsistent access to shared data

Example of Thread Usage

A multithreaded Web server.



Simplified code...

- (Note: Does not show the synchronization primitives)

```
while (TRUE) {
    get_next_request(&buf);
    handoff_work(&buf);
}
```

(a)

“dispatcher” thread

```
while (TRUE) {
    wait_for_work(&buf)
    look_for_page_in_cache(&buf, &page);
    if (page_not_in_cache(&page))
        read_page_from_disk(&buf, &page);
    return_page(&page);
}
```

(b)

“worker” thread

Thread Advantages

- Performance: interactions and switches among threads is faster than for processes
- Even when executed on a uniprocessor, threads are often seen as a more “natural” way to program certain classes of applications:
 - Servers. How so? Can be scheduled to run
 - Simulation. How so? in multiprocessors
 - Pipelined data processing (producer/consumer)
- Kernel-level threads can “automatically” take advantage of a multi-processor/core system.

Parallel Speedup

- Speedup for P processors:
 - $S_N = T_1/T_N$, where T_1 is serial time, T_N is time on N processors
- Efficiency, $E_N = S_N/P = T_1/(P*T_N)$
- Example:
 - $T_1 = 12\text{s}$, $T_4 = 4\text{s}$
 - $S_4 = ?$ 12/4 = 3
 - $E_4 = ?$ 12/(4*4) = 12/16 = 0.75

Amdahl's Law

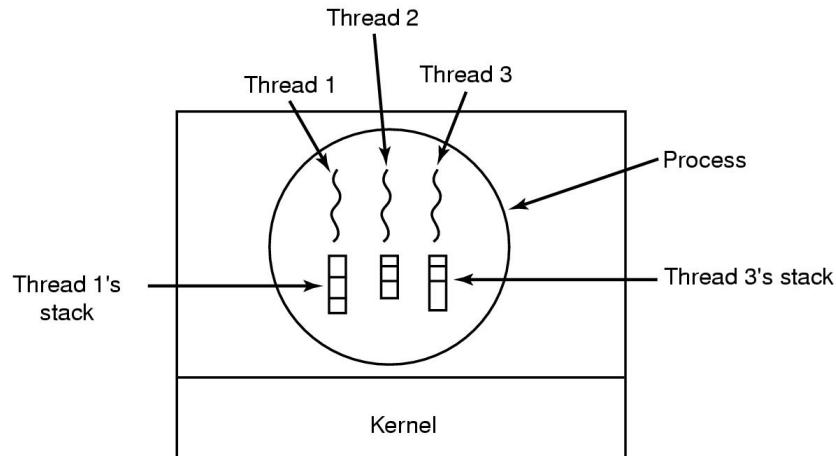
- Limit on speedup depends on the fraction of code that can (and cannot) be parallelized
- $S_N = T_1/T_N < 1/((1-f) + f/N)$, where
 - f is the fraction of code that can be parallelized
- Example: $f = 0.9$
 - $S_4 < 1/(0.1+0.9/4) = 3.1$
 - $S_{10} < 1/(0.1+0.9/10) = 5.3$
 - $S_{100} < 1/(0.1+0.9/100) = 9.2$
 - $S_{1000} < 1/(0.1+0.9/1000) = 9.9$

Multiprocessor Thread Performance

- Thread-based applications may be executed on either uniprocessors or multiprocessors
- On a multiprocessor, performance depends on the level of OS support:
 - User-level threads? $\text{max speed: } 1.0$
 - Kernel-level threads ≥ 1.0

Each Thread Needs...

- Whether implemented as kernel- or user-level threads
 - Clearly, each thread needs its own stack.
 - Per-thread global variables are highly desirable.



CSE 410 2-3 Threads

13

Thread-Specific and Shared Resources

- Each thread has its own context:
 - program counter, stack, register set, etc.
- In addition to address space, all threads within a single process typically share:
 - open files, signals and other kernel services

Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

CSE 410 2-3 Threads

13

Thread/Process Issues

- Semantics of the interaction with the kernel can complicate programming
 - Has been a major issue in Linux over the years...
- Examples:
 - When a thread invokes fork(), are all threads instantiated in the new process?
 - What if one thread is blocked waiting for input from a keyboard? Are other threads now blocked? What if two threads are waiting? Which thread gets the input?
 - What if one thread closes a file while another is still reading from it?
 - And so on...

A few more details on User-Level Threads

- Thread table contains entry for each thread:
 - Program ctr, stack ptr, state, thread context, etc.
- Run-time system (like a small operating system) manages threads within the process
 - Implements communication, synchronization among threads
 - Schedules and dispatches threads (thread switching)
- Advantages over process-based design, or even kernel-supported threads?
no kernel support required.
interaction/sync faster because all you do is function calls.
thread creation.

Problems with User-Level Threads

- Blocking system calls
 - One blocked thread will block all threads in process
 - Possible solution: Wrap such calls with code to check in advance whether call will block. How? [Select system call select\(\)](#)
- Other blocking situation?
[I/O exceptions, page fault](#)
- How to implement time slicing?
[sigalarm](#)
- Problems arise for I/O bound applications
 - System call overhead subsumes cost of switching threads
 - Why not just support threads in kernel...

Kernel-supported threads

- Kernel maintains context of each thread within a process
- Thread creation, suspension, synchronization, are implemented with system calls to the kernel.
- Threads **scheduled** independently!
[slower creation time.](#)
- Solves the blocking problem. Disadvantages?
[interaction](#)
- Issues difficult with either type thread: some library routines may not be reentrant, handling per process entities such as signals, open files...
- ROOT of the problem???
[processes exist before multiple threads.](#)
[threads added after the fact.](#)

Threads in Linux

- Like any other operating system, Linux supports user-level threads
- With the clone() system call, also supports kernel-level threads
 - Creates child process (with own task_struct) in same address space of parent
 - Used to partially support POSIX threads
- In early 2000s, Native POSIX Thread Library (NPTL) developed to fully support threads
- POSIX threads, or **Pthreads**

Example: POSIX Threads

- To support thread portability and uniformity, at least among Unix versions, IEEE defined the POSIX thread standard (**Pthreads**)
- Standardizes 60 function calls and various thread attributes and semantics
- Examples:

Thread call	Description
Pthread_create	Create a new thread
Pthread_exit	Terminate the calling thread
Pthread_join	Wait for a specific thread to exit
Pthread_yield	Release the CPU to let another thread run
Pthread_attr_init	Create and initialize a thread's attribute structure
Pthread_attr_destroy	Remove a thread's attribute structure

Simple Pthreads Program

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

#define NUMBER_OF_THREADS 10

void *print_hello_world(void *tid)
{
    /* This function prints the thread's identifier and then exits. */
    printf("Hello World. Greetings from thread %d0, tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[])
{
    /* The main program creates 10 threads and then exits. */
    pthread_t threads[NUMBER_OF_THREADS];
    int status, i;

    for(i=0; i < NUMBER_OF_THREADS; i++) {
        printf("Main here. Creating thread %d0, i);
        status = pthread_create(&threads[i], NULL, print_hello_world, (void *)i);
        ...
        if (status != 0) {
            printf("Oops. pthread_create returned error code %d0, status);
            exit(-1);
        }
    }
    exit(NULL);
}
```

CSE 410 2-3 Threads

22

Other Interfaces

- Threads are supported in various programming languages and libraries, e.g.,
 - C++11, std::thread
 - Java Thread Object
 - Boost C++ Library, boost::thread
- In most cases, the programming primitives simply “wrap” a lower level implementation such as Pthreads
 - Advantage? [Schedulable](#)
- In CSE 410, we'll focus on Pthreads
 - Differences are largely just syntax issues...

Last comment: Multithreaded Kernel

- Supporting kernel reentrancy and even kernel-level threads on a uniprocessor is one thing...
 - Don't have to worry about preemption by other threads
 - Locks around data structures manipulated by interrupt handlers
 - Longer term locks (waiting for some resource) built into the logic of the kernel (sleep and wakeup)
- Implementing a multithreaded kernel is quite another.
- Supports concurrent execution of kernel code on multiple processors
- Approach: Lock everything!

Brief Introduction to Pthreads

(examples in C)

Pthreads

- Pthreads is a POSIX standard for describing a thread model, it specifies the API and the semantics of the calls.
- This model is popular – nowadays practically all major thread libraries on Unix-like systems are pthreads-compatible
- In Linux, pthreads are realized as **kernel-level** threads.
 - Implications?

Process vs. Threads

- Threads share:
 - process instructions (text)
 - global data
 - open files (descriptors)
 - signals and signal handlers
 - current working directory
 - user and group id
- Unique to thread
 - thread ID
 - register context
 - stack
 - signal mask
 - priority, return value

Compiling pthread-based programs

- Include `pthread.h` in the main file
- Compile program with `-lpthread` or `-pthread`
 - `gcc -o test test.c -lpthread`
 - `g++ -pthread test.cpp`

Thread-related calls

- Types: `pthread_t` – type of a thread
- Some calls:

```
int pthread_create(pthread_t *thread,  
                    const pthread_attr_t *attr,  
                    void * (*start_routine)(void *),  
                    void *arg);  
  
int pthread_join(pthread_t thread, void  
                 **status);  
  
int pthread_detach();  
void pthread_exit();
```

As with other primitives, always check return values on pthread functions

Thread features

- No explicit parent/child model, except main thread holds process info
- To enable other threads to continue execution after main thread is finished, need to call `pthread_exit` in main, don't just fall through
- Individual threads can call `pthread_exit` or just fall through (implicit call)
- `pthread_join` used to wait for another thread to exit
- Detached threads are those which cannot be joined (can also set this at thread creation)

Thread Creation and Termination

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *print_message_function( void *ptr );

main()
{
    pthread_t thread1, thread2;
    const char *message1 = "Thread 1";
    const char *message2 = "Thread 2";
    int iret1, iret2;

    /* Create independent threads each of which will execute function */

    iret1 = pthread_create( &thread1, NULL, print_message_function, (void*) message1 );
    iret2 = pthread_create( &thread2, NULL, print_message_function, (void*) message2 );

    /* Wait till threads are complete before main continues. Unless we */
    /* wait we run the risk of executing an exit which will terminate */
    /* the process and all threads before the threads have completed. */

    pthread_join( thread1, NULL );
    pthread_join( thread2, NULL );

    printf("Thread 1 returns: %d\n",iret1);
    printf("Thread 2 returns: %d\n",iret2);
    exit(0);
}

void *print_message_function( void *ptr )
{
    char *message;
    message = (char *) ptr;
    printf("%s \n", message);
}
```

Mutexes

- Used to
 - control access to shared resources
 - prevent race conditions
- Pthread mutexes work only among threads, not among processes
- Primitives:
 - `pthread_mutex_t mutex_variable = PTHREAD_MUTEX_INITIALIZER;`
 - `pthread_mutex_lock(&mutex_variable);`
 - `pthread_mutex_unlock(&mutex_variable);`
 - `pthread_mutex_trylock()` – tries but does not block...

CSE 410 2-3 Pthreads

74

Simple Mutex Example

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *functionC();
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;
int counter = 0;

main()
{
    int rc1, rc2;
    pthread_t thread1, thread2;

    /* Create independent threads each of which will execute functionC */

    if( (rc1=pthread_create( &thread1, NULL, &functionC, NULL)) )
    {
        printf("Thread creation failed: %d\n", rc1);
    }

    if( (rc2=pthread_create( &thread2, NULL, &functionC, NULL)) )
    {
        printf("Thread creation failed: %d\n", rc2);
    }

    /* Wait till threads are complete before main continues. Unless we */
    /* wait we run the risk of executing an exit which will terminate */
    /* the process and all threads before the threads have completed. */

    pthread_join( thread1, NULL);
    pthread_join( thread2, NULL);

    exit(0);
}

void *functionC()
{
    pthread_mutex_lock( &mutex1 );
    counter++;
    printf("Counter value: %d\n", counter);
    pthread_mutex_unlock( &mutex1 );
}
```

CSE 410 2-3 Pthreads

75

Join

- Wait for a thread to finish
- Example: wait for a collection of worker threads to finish their tasks
 - `pthread_join()` - wait for termination of another thread
 - `pthread_self()` - return identifier of current thread

Join Example

```
#include <stdio.h>
#include <pthread.h>

#define NTHREADS 10
void *thread_function(void *);  
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;  
int counter = 0;

main()
{
    pthread_t thread_id[NTHREADS];
    int i, j;

    for(i=0; i < NTHREADS; i++)
    {
        pthread_create( &thread_id[i], NULL, thread_function, NULL );
    }

    for(j=0; j < NTHREADS; j++)
    {
        pthread_join( thread_id[j], NULL );
    }

    /* Now that all threads are complete I can print the final result. */
    /* Without the join I could be printing a value before all the threads */
    /* have been completed. */
    printf("Final counter value: %d\n", counter);
}

void *thread_function(void *dummyPtr)
{
    printf("Thread number %ld\n", pthread_self());
    pthread_mutex_lock( &mutex1 );
    counter++;
    pthread_mutex_unlock( &mutex1 );
}
```

Thread Programming Problems

Logical execution is nondeterministic

- Race conditions (BTW, what is a race condition?)
 - Cannot predict how kernel schedules threads
 - Use mutexes and condition variables to help control it
- Code that is not thread-safe (e.g., reentrant)
 - Code was written to be linked to and used by a process
 - Called routines should not contain static or global variables unless they are protected by mutexes
 - Can protect the entire call with a mutex
- Mutex deadlock
 - Mutex is applied but not unlocked
 - Especially a problem when multiple mutexes are involved

Simple Example

- Grab one mutex
- If waiting for another, free first

```
...
pthread_mutex_lock(&mutex_1);
while ( pthread_mutex_trylock(&mutex_2) ) /* Test if already locked */
{
    pthread_mutex_unlock(&mutex_1); /* Free resource to avoid deadlock */
    ...
    /* stall here */
    ...
    pthread_mutex_lock(&mutex_1);
}
count++;
pthread_mutex_unlock(&mutex_1);
pthread_mutex_unlock(&mutex_2);
...
```

Deadlock Example: Order is important

```
void *function1()
{
    ...
    pthread_mutex_lock(&lock1);           // Execution step 1
    pthread_mutex_lock(&lock2);           // Execution step 2 DEADLOCK!!!
    ...
    ...
    pthread_mutex_lock(&lock2);
    pthread_mutex_lock(&lock1);
    ...
}

void *function2()
{
    ...
    pthread_mutex_lock(&lock2);           // Execution step 2
    pthread_mutex_lock(&lock1);
    ...
    ...
    pthread_mutex_lock(&lock1);
    pthread_mutex_lock(&lock2);
    ...
}

main()
{
    ...
    pthread_create(&thread1, NULL, function1, NULL);
    pthread_create(&thread2, NULL, function2, NULL);
    ...
}
```

CSE 410 2-3 Pthreads

86

Summary

- What we have learned
 - Basic concept of multiple threads executing within one process
 - User-level vs. kernel-level threads
 - Contextual information needed for each thread
 - Speedup and Amdahl's Law
 - Semantics issues due to process model
 - Brief introduction to pthreads programming
 - Example usage of simple concurrency control primitives
- Next up: More complete treatment of concurrency (“meaty” topic applicable not only to threads!)

CSE 410 2-3 Pthreads

87