
Concurrency

Mutual Exclusion and Synchronization

Outline

- Read Stallings, 5.1, 5.2, 5.3, 5.4, 5.7
- Topics:
 - Critical sections and mutual exclusion
 - Pure software solutions
 - Hardware-based solutions
 - Semaphores
 - Classic Concurrency Problems
 - Producer consumer
 - Dining philosophers
 - Readers/writers problem

Problems with Concurrent Execution

- Concurrent processes (or threads) often need to share data (maintained either in shared memory or files) and resources
- If there is no controlled access to shared data, some processes might obtain an inconsistent view of this data
- The action performed by concurrent processes will then depend on the order in which their execution is interleaved, which typically is **not deterministic**

A Simple Example

- Process P1 and P2 are running this same procedure and have access to the same variable “a”
- Processes can be interrupted anywhere
- If P1 is first interrupted after user input and P2 executes entirely
- Then the character echoed by P1 will be the one read by P2
!!

```
static char a;

void echo()
{
    cin >> a;
    cout << a;
}
```

The Critical Section Problem

- When a process executes code that manipulates shared data (or resource), we say that the process is in a **critical section** (CS) for that shared data
- Sometimes referred to as a *critical region*
- The execution of critical sections must be **mutually exclusive**: At any time, only one process is allowed to execute in its critical section (even with multiple CPUs)
- Each process must **request** the permission to enter the critical section (CS)

Critical Sections

- The section of code implementing this request is called the **entry section**
- The critical section (CS) might be followed by an **exit section**
- The **critical section problem** is to design a protocol that the processes can use so that their action will not depend on the order in which their execution is interleaved (possibly on many processors)

Mutual Exclusion and OS

- Operating systems are usually responsible for providing processes with mutually exclusive access to system resources
 - relatively easy on a uniprocessor. Why? *Only one process in kernel at a time*
 - for a multiprocessor, kernel will need to implement critical sections in its own code (multiple processes in kernel)
- Critical sections sometimes need to be implemented by user programs, even on a uniprocessor. Example? *Same file, share memory*
- Obviously, SMP kernels implement critical sections everywhere. Even uniprocessor kernels need critical sections. Why? *Interrupt Handlers*

CSE 410 2-4 Concurrency

Framework for analysis of solutions

- Each process executes at nonzero speed but no assumption on the relative speed of n processes
- General structure of a process:

repeat
 entry section
 critical section
 exit section
 remainder section
forever
- Many CPUs may be present but memory hardware prevents simultaneous access to the same memory location
- No assumption about order of interleaved execution
- For solutions: we need to specify entry and exit sections

Requirements for a valid solution

- Mutual Exclusion
 - At any time, at most one process can be in the critical section (CS)
- Progress
 - If no process is executing in its CS while some processes wish to enter, any process that requests entry should be granted such without delay
- Processes remain in critical section for finite time
- A waiting process cannot be delayed indefinitely (no deadlock or starvation)

Types of Solutions

- “Pure” software solutions
 - algorithms where correctness does not rely on any other assumptions beyond framework described earlier
 - Simply implemented in code, without need for atomic instructions
- Hardware solutions
 - rely on special machine instructions
 - e.g., test-and-set, xchg, **cmpxchg**
- Operating system solutions
 - provide system calls and data structures to the user programs; mutual exclusion implemented by kernel

Pure Software Solutions

- We consider the case of two processes
 - We assess three different algorithms
- Approach generalizes to more processes
- Notation
 - We start with 2 processes: P0 and P1
 - When presenting process P_i , P_j always denotes the other process ($i \neq j$)

Algorithm 1

- The shared variable `turn` is initialized (to 0 or 1) before executing any P_i
- P_i 's critical section is executed iff ($\text{turn} == i$)
- P_i is **busy waiting** if P_j is in CS: mutual exclusion is satisfied
- Problem?

```
Process  $P_i$ :  
repeat  
    while ( $\text{turn} \neq i$ ) {};  
    CS  
     $\text{turn} = j$ ;  
    RS  
forever
```

Alternations prohibits progress

Algorithm 2

- Add a boolean variable for each process: flag[0] and flag[1]
- P_i signals that it is ready to enter its CS by: flag[i] = true
- Mutual Exclusion is satisfied
- Problem?

```
Process  $P_i$ :  
repeat  
    flag[i] = true;  
    while(flag[j]){};  
    CS  
    flag[i] = false;  
    RS  
forever
```

Algorithm 3 (Peterson's Algorithm)

- Initialization:
flag[0] = flag[1] = false
turn = 0 or 1
- Willingness to enter CS
specified by setting
flag[i] = true
- If both processes attempt
to enter their CS
simultaneously, only one
turn value will prevail
- Exit section: specifies that
 P_i is unwilling to enter CS

```
Process  $P_i$ :  
repeat  
    flag[i] = true;  
    turn = j;  
    do {} while  
        (flag[j] and (turn==j));  
    CS  
    flag[i] := false;  
    RS  
forever
```

Algorithm 3: Proof of Correctness

- Mutual exclusion:
- Proof by contradiction:
- Assume both P0 and P1 are in their CS
 - then $\text{flag}[0] = \text{flag}[1] = \text{true}$
 - but the test for entry cannot have been true for both processes at the same time (because turn favors one);
 - therefore one process must have entered its CS first (without loss of generality, say P0)
 - but this means that P1 could not have found $\text{turn} = 1$ and therefore could not have entered its CS (i.e. contradiction)

Algorithm 3: Progress

Consider P0 blocked in the entry loop:

- Case I: (Stuck)
 - P1 is not interested in entering its CS
 - then $\text{flag}[1] = \text{false}$
 - hence the while loop is false for P0 and it can go
- Case II: (Deadlock)
 - P1 is also blocked at the while loop
 - impossible, because $\text{turn} = 0$ or 1
 - hence the while loop is false for some process and it can go

Proof: Bounded Waiting

- Case III: (Starvation)
- P1 is executing its CS repeatedly
 - upon exiting its CS, P1 sets $\text{flag}[1] = \text{false}$
 - hence the while loop is false for P0 and it can go (sufficient?)
- However, P1 may attempt to re-enter its CS before P0 has a chance to run.
 - but to re-enter, P1 sets $\text{flag}[1]$ to true and sets turn to 0
 - hence the while loop is true for P1 and it waits
 - the while loop is now false for P0 and it can go
- Note: Can extend to N processes (Bakery Algorithm)

What about process failures?

- If all three criteria (mutex, progress, bounded waiting) are satisfied, then a valid solution will provide robustness against failure of a process in its remainder section (RS)
 - since failure in RS is just like having an infinitely long RS
- However, no valid solution can provide robustness against a process failing in its critical section (CS) – other recovery methods are needed

Drawbacks of software solutions

- Processes that are requesting to enter in their critical section are **busy waiting** (consuming processor time needlessly)
- Can a process busy-wait for its entire timeslice?
- If critical sections are long, it would be more efficient to **block** (i.e., suspend) those processes that are waiting...
- BTW, when is busy waiting considered acceptable?

Kernel locks

Hardware Solutions: Interrupt Disabling

- Not allowed for user processes!
- What about the kernel?
On a uniprocessor:
- On a multiprocessor:
Atomic instructions

```
Process Pi:  
repeat  
    disable interrupts  
    critical section  
    enable interrupts  
    remainder section  
forever
```

Special Machine Instructions

- Processor designers have implemented machine instructions that perform two actions **atomically (indivisible)** on the same memory location (e.g., reading and writing)
- The execution of such an instruction is **mutually exclusive** (even with multiple CPUs)
- Such instructions can be used to provide mutual exclusion
- To satisfy all three requirements of the CS problem (incl. bounded waiting) requires additional mechanisms.

test-and-set semantics

- A C++ description of test-and-set:

```
bool testset(int& i)
{
    if (i==0) {
        i=1;
        return true;
    } else {
        return false;
    }
}
```

- An algorithm that uses testset for Mutual Exclusion:
- Shared variable b is initialized to 0
- Only the first P_i who sets b enters the CS

```
Process Pi:
repeat
    repeat{}
    until testset(b) ;
    CS
    b:=0;
    RS
forever
```

test-and-set properties

- Mutual exclusion is preserved: if P_i enter CS, the other P_j are **busy waiting**
- Problem: still uses busy waiting
- When P_i exit CS, the selection of the P_j who will enter CS is arbitrary: **no bounded waiting**. Hence **starvation** is possible
- Other instructions on various processors, such as `xchg(a,b)` swaps the content of `a` and `b`.

23
CSE 410 2-4 Concurrency

Using xchg for mutual exclusion

- Shared variable `b` is initialized to 0
- Each P_i has a local variable `k`
- The only P_i that can enter CS is the one who finds `b=0`
- This P_i excludes all the other P_j by setting `b` to 1

```
Process  $P_i$ :  
repeat  
    k:=1  
    repeat xchg(k,b)  
    until k=0;  
        CS  
    b:=0;  
        RS  
forever
```

But `xchg` suffers from the same drawback as test-and-set

OS-Based Solutions

- We saw that we can implement mutual exclusion:
 - Purely in code, with no help from hardware or the OS
 - By disabling interrupts (but only for uniprocessor kernel)
 - With atomic instructions (either user or kernel)
 - Problem with these approaches?
- Alternative approach for supporting user processes: Implement mutual exclusion (and process synchronization) in the OS.
- Examples:
 - mutexes, semaphores, condition variables, monitors

Advantage of OS support?

Semaphores

- Synchronization tool (usually provided by the OS) that does not require busy waiting
- A semaphore S is an integer variable that, apart from initialization, can only be accessed via **atomic and mutually exclusive** operations:
 - wait(S) -- sometimes P(S)
 - signal(S) -- sometimes V(S)
- To avoid busy waiting: when a process has to wait, it will be put in a **queue** of blocked processes waiting for the same event

Semaphores

- Hence, a semaphore can be implemented as a structure with two fields:
 - count: integer
 - queue: list of processes
- When a process must wait for a semaphore S, it is blocked and put on the semaphore's queue
- The signal operation removes (according to a fair policy like FIFO) one process from the queue and puts it in the list of ready processes

27

Semaphore's operations (atomic)

```
wait(S) {  
    S.count--;  
    if (S.count < 0) {  
        block this process in S.queue  
    }  
}
```

```
signal(S) {  
    S.count++;  
    if (S.count <= 0) {  
        remove a process P from S.queue  
        place this process P on ready list  
    }  
}
```

S.count must be initialized to a nonnegative value (depending on application)

Semaphores: Observations

- When $S.count \geq 0$: the number of processes that can execute `wait(S)` without being blocked = $S.count$
- When $S.count < 0$: the number of processes waiting on S is $= |S.count|$
- Atomicity and mutual exclusion: no two process can be in `wait(S)` and `signal(S)` (on the same semaphore, S) at the same time (even with multiple CPUs)
- Hence the blocks of code defining `wait(S)` and `signal(S)` are, in fact, critical sections

Using semaphores for critical sections!

- For n processes
- Initialize $S.count$ to 1
- Then only 1 process is allowed into CS (mutual exclusion)
- To allow k processes into CS, we initialize $S.count$ to k

```
Process  $P_i$ :  
repeat  
    wait(S) ;  
    CS  
    signal(S) ;  
    RS  
forever
```

Using semaphores for synchronization

- We have two processes: P1 and P2
- Statement S1 in P1 needs to be performed before statement S2 in P2
- Then define a semaphore “synch”
- Initialize synch to 0
- Proper synchronization is achieved by having in P1:
 - S1;
 - signal(synch);
- And having in P2:
 - wait(synch);
 - S2;

Semaphores in Linux

- Included in POSIX standard
- Can be used by processes and threads
- Implementation:
 - integer whose value is never allowed to fall below zero
 - sem_post(3): increment the semaphore value by one
 - sem_wait(3): decrement the semaphore value by one. If the value of a semaphore is currently zero, then a sem_wait operation will block until the value becomes greater than zero.


The producer/consumer problem

- A **producer process** produces information that is consumed by a **consumer process**
 - Ex1: a print program produces characters that are consumed by a printer
 - Ex2: an assembler produces object modules that are consumed by a loader
- We need a **buffer** to hold items that are produced and eventually consumed
- A common paradigm for cooperating processes

P/C: unbounded buffer

- We assume first an **unbounded** buffer consisting of a linear array of elements
- `in` points to the next item to be produced
- `out` points to the next item to be consumed

shaded area indicates portion of buffer that is occupied



P/C: unbounded buffer

- We need a semaphore S to perform mutual exclusion on the buffer: only 1 process at a time can access the buffer
- We need another semaphore N to synchronize producer and consumer on the number N of items in the buffer ($N = \text{in} - \text{out}$)
 - Note: an item can be consumed only after it has been created

P/C: unbounded buffer

- The producer is free to add an item into the buffer at any time (since buffer is unbounded)
 - it performs $\text{wait}(S)$ before appending and $\text{signal}(S)$ afterwards to prevent customer access
 - It also performs $\text{signal}(N)$ after each append to increment N
- The consumer must first $\text{wait}(N)$ to see if there is an item to consume and then use $\text{wait}(S)/\text{signal}(S)$ to access the buffer

Solution of P/C: unbounded buffer

```
Initialization:
    S.count:=1;
    N.count:=0;
    in:=out:=0;

append(v) :
    b[in]:=v;
    in++;

take() :
    w:=b[out];
    out++;
    return w;

Producer:
repeat
    produce v;
    wait(S);
    append(v);
    signal(S);
    signal(N);
forever

Consumer:
repeat
    wait(N);
    wait(S);
    w:=take();
    signal(S);
    consume(w);
forever


■ critical sections
```

38

P/C: unbounded buffer

- Remarks:
 - Putting signal(N) inside the CS of the producer (instead of outside) has no effect since the consumer must always wait for both semaphores before proceeding
 - The consumer must perform wait(N) before wait(S), otherwise **deadlock** occurs if consumer enter CS while the buffer is empty
- BTW, using semaphores and debugging solutions is not always trivial...

P/C: **finite** circular buffer of size k



- can consume **only** when number N of (consumable) items is at least 1
- can produce **only** when number E of empty spaces is at least 1

P/C: finite circular buffer of size k

- As before:
 - we need a semaphore S to have mutual exclusion on buffer access
 - we need a semaphore N to synchronize producer and consumer on the number of consumable items
- **In addition:**
 - we need a semaphore E to synchronize producer and consumer on the number of **empty** spaces, so producer knows if it can append

Solution

Initialization: S.count:=1; in:=0;
N.count:=0; out:=0;
E.count:=k;

append(v) :

b[in]:=v;

in:=(in+1)mod k;

take() :

w:=b[out];

out:=(out+1)mod k;

return w;

Producer:

repeat

produce v;

wait(E);

wait(S);

append(v);

signal(S);

signal(N);

forever

Consumer:

repeat

wait(N);

wait(S);

w:=take();

signal(S);

signal(E);

consume(w);


forever

■ critical sections

42

The Dining Philosophers Problem

- Classic synchronization problem...
- Five philosophers who only eat and think
- Each need to use two forks for eating
- We have only 5 forks
- Illustrates the difficulty of allocating resources among process without deadlock and starvation



The Dining Philosophers Problem

- Each philosopher is a process
- One semaphore per fork:
 - fork: array[0..4] of semaphores
 - Initialization: fork[i].count:=1 for i:=0..4
- Solution is shown a the right.
- Problem?

```
Process Pi:
repeat
  think;
  wait(fork[i]);
  wait(fork[i+1 mod 5]);
  eat;
  signal(fork[i+1 mod 5]);
  signal(fork[i]);
forever
```

The Dining Philosophers Problem

- One solution: admit only four philosophers at a time.
- Then one philosopher can always eat when the other three are holding one fork
- Hence, we can use another semaphore T that would limit at four the number of philosophers “sitting at the table”
- Initialize: T.count:=4

```
Process Pi:
repeat
  think;
  wait(T);
  wait(fork[i]);
  wait(fork[i+1 mod 5]);
  eat;
  signal(fork[i+1 mod 5]);
  signal(fork[i]);
  signal(T);
forever
```

Deadlock, livelock, starvation

- Deadlock: permanent blocking of a set of processes competing for resources and/or communicating
- Three conditions for deadlock:
 - **mutual exclusion** in access to resources
 - a process may **hold** resources **and wait** for others
 - no preemption – cannot forcibly remove resource from process holding it.
- Livelock – set of processes continuously change states (regarding resource acquisition) but do not make progress.

Examples

Readers/Writers Problem

- Classic problem in computer science
- A data area is shared among many processes
 - some processes only read the data area, (readers) and some only write to the data area (writers)
- Conditions that must be satisfied:
 1. any number of readers may simultaneously read the file
 2. only one writer at a time may write to the file
 3. if a writer is writing to the file, no reader may read it

Readers have priority

- Semaphore mutex
 - Simply protects readcount var
 - Could just be a mutex, if such primitives are available
- Semaphore wsem
 - For mutually exclusive access to critical section
- Problem with this solution?

```
semaphore wsem=1, mutex=1; readcount=0;
```

```
writer()
{
    wait(wsem);
    // Do the writing (Critical Section)
    signal(wsem);
}
```

```
reader()
{
    wait(mutex);
    readcount++;
    if (readcount == 1)
        wait(wsem);
    signal(mutex);
    // Do the reading (Critical Section)
    wait(mutex);
    readcount--;
    if (readcount == 0)
        signal(wsem);
    signal(mutex);
}
```


Writers have priority

- Somewhat more complicated
- Need to keep track of writers who might be waiting
- Need to enable them to “jump ahead” of waiting readers
- Some additional variables:

```
// global variables. this solution uses only semaphores, not mutexes
int readcount = 0; // how many readers are currently in
int writecount = 0; // how many writers are currently in (or want in)
semaphore readcount_mutex = 1; // mutually exclusive updating of readcount
semaphore writecount_mutex = 1; // mutually exclusive updating of writecount
semaphore reader_gate = 1; // why needed?
semaphore write_sem = 1; // mutual exclusion on critical section
semaphore read_sem = 1; // readers have to wait if any writer desires access
```

Writers have priority (writer code)

```
writer()
{
    // indicate writer is waiting
    wait(writecount_mutex);
    writecount++;
    if (writecount == 1)
        wait(read_sem);
    signal(writecount_mutex);
    wait(write_sem);

    // writing is performed here

    signal(write_sem);

    wait(writecount_mutex);
    writecount := writecount - 1;
    if (writecount == 0) // no other writers are waiting, so let in any waiting readers
        signal(read_sem);
    signal(writecount_mutex);
}
```

Writers have priority (reader code)

```
reader()
{
    wait(reader_gate); // first gate for readers
    wait(read_sem); // is it ok for reader to wait
    wait(readcount_mutex); // update readcount
    readcount := readcount + 1;
    if (readcount == 1)
        wait(write_sem); //
    signal(readcount_mutex);
    signal(read_sem);
    signal(reader_gate);

    //reading is performed here

    // reading is done; decrement readcount; if all readers done, release write_sem
    wait(readcount_mutex);
    readcount = readcount - 1;
    if (readcount == 0)
        signal(write_sem); // ok to let someone in (either reader or writer), writers priority
    signal(readcount_mutex);
}
```

Example

reader_gate:




read_sem:



write_sem:



Shared Resource



Example

reader_gate:




read_sem:



write_sem:



Shared Resource



Example

reader_gate:




read_sem:



write_sem:



Shared Resource



Summary

- In this chapter we have addressed the issue of concurrent execution among processes/threads
- Key issues
 - Mutual exclusion (progress, no starvation, no deadlock)
 - Synchronization
- Methods
 - Mutexes, semaphores (others: monitors, condition vars)
- Classic problems that arise in OS and general CS:
 - Resource dependencies (dining philosophers)
 - Producer/consumer problem
 - Readers/writers problem